

Program 7

```
In [1]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.neighbors import KNeighborsClassifier

df=pd.read_csv(r"D:\PRIYA\SEM 05\dataset\iris.csv")
df.head()
```

```
Out[1]:   sepal_length  sepal_width  petal_length  petal_width  species
0           5.1         3.5          1.4         0.2    setosa
1           4.9         3.0          1.4         0.2    setosa
2           4.7         3.2          1.3         0.2    setosa
3           4.6         3.1          1.5         0.2    setosa
4           5.0         3.6          1.4         0.2    setosa
```

```
In [2]: x=df.drop(['species'],axis=1)
y=df['species']
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=42)

Scaler=StandardScaler()
x_train=Scaler.fit_transform(x_train)
x_test=Scaler.transform(x_test)

classifier=KNeighborsClassifier(n_neighbors=5)
classifier.fit(x_train, y_train)
y_pred = classifier.predict(x_test)
```

```
In [6]: accuracy = accuracy_score(y_test,y_pred)
print(f"Accuracy: {accuracy:.2f}\n")
print("Confusion Matrix:")
print(confusion_matrix(y_test,y_pred))
print("Classification Report:")
print(classification_report(y_test,y_pred))

Accuracy: 1.00

Confusion Matrix:
[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]
Classification Report:
      precision    recall  f1-score   support

        setosa       1.00     1.00     1.00      19
versicolor       1.00     1.00     1.00      13
 virginica       1.00     1.00     1.00      13

   accuracy         -         -         -      45
    macro avg       1.00     1.00     1.00      45
 weighted avg     1.00     1.00     1.00      45
```

```
In [7]: for true,pred in zip (y_test,y_pred):
    if true==pred:
        print(f"correct:true label={true}, predicted label={pred}")
    else:
        print("false:true label={true}, predicted label={pred}")
```

Program 8

```
In [8]: import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, precision_score, recall_score, f1_score

data=pd.read_csv(r"D:\PRIYA\SEM 05\dataset\NaiveText.csv")
#data=pd.read_csv(r'/content/IMDB Dataset.csv')

data.head()
```

```
Out[8]:      message  label
0   I love this sandwich    1
1 This is an amazing place    1
2 I feel very good about these beers    1
3 This is my best work    1
4 What an awesome view    1
```

```
In [10]: x=data['message']
y=data['label']

#x=data['review']
#y=data['sentiment']

x_train,x_test, y_train,y_test= train_test_split(x,y,test_size=0.2,random_state=42)

vectorizer = CountVectorizer()
x_train_vectorized = vectorizer.fit_transform(x_train)
x_test_vectorized = vectorizer.transform(x_test)

classifier = MultinomialNB()
classifier.fit(x_train_vectorized, y_train)
y_pred = classifier.predict(x_test_vectorized)
print(y_pred)
```

```
\nAccuracy: 1.0
Precision: 1.0
Recall: 1.0
F1 Score: 1.0
Confusion Matrix
```

```
[[2 0]
 [0 2]]
Classification Report
      precision    recall  f1-score   support
          0       1.00     1.00    1.00      2
          1       1.00     1.00    1.00      2

   accuracy                           1.00      4
  macro avg       1.00     1.00    1.00      4
weighted avg       1.00     1.00    1.00      4
```

Program 9

Program 10

```
In [11]: import pandas as pd
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import LabelEncoder

data = pd.read_csv(r"D:\PRIYA\SEM 05\dataset\PlayTennis.csv")
data.head()
```

```
Out[11]:   Outlook Temperature Humidity Wind Play Tennis
0      Sunny        Hot     High  Weak     No
1      Sunny        Hot     High Strong    No
2    Overcast       Hot     High  Weak    Yes
3       Rain       Mild     High  Weak    Yes
4       Rain      Cool  Normal  Weak    Yes
```

```
In [12]: le = LabelEncoder()
categorical_cols = ['Outlook', 'Temperature', 'Humidity', 'Wind', 'Play Tennis']

for col in categorical_cols:
    data[col] = le.fit_transform(data[col])
data.head(5)

X = data.drop('Play Tennis', axis=1)
y = data['Play Tennis']
x_train,x_test,y_train,y_test = train_test_split(X,y,test_size=0.33,random_state=42)

clf = GaussianNB()
clf.fit(x_train,y_train)
ypred = clf.predict(x_test)
print(ypred)

accuracy = accuracy_score(y_test, ypred)
print("Accuracy:", accuracy)
```

[1 0 0 1 1]
Accuracy: 0.6

Program 11


```
In [23]: import numpy as np
import pandas as pd
from sklearn.preprocessing import LabelEncoder,StandardScaler
from sklearn.cluster import KMeans,DBSCAN
from sklearn.metrics import silhouette_score
import matplotlib.pyplot as plt
import seaborn as sns
df=pd.read_csv(r"D:\PRIYA\SEM 05\dataset\iris.csv")
df.head(5)
```

```
Out[23]:   sepal_length  sepal_width  petal_length  petal_width  species
0           5.1         3.5          1.4         0.2    setosa
1           4.9         3.0          1.4         0.2    setosa
2           4.7         3.2          1.3         0.2    setosa
3           4.6         3.1          1.5         0.2    setosa
4           5.0         3.6          1.4         0.2    setosa
```

```
In [52]: label_encoder=LabelEncoder()
df["species"] = label_encoder.fit_transform(df["species"])
features=['sepal_length','sepal_width','petal_length','petal_width']
x=df[features]
scaler=StandardScaler()
x_scaled=scaler.fit_transform(x)
wcss=[]
for i in range(1,11):
    km=KMeans(n_clusters=i)
    km.fit(x_scaled)
    wcss.append(km.inertia_)

plt.figure(figsize=(4,3))
plt.plot(range(1,11),wcss,marker='o',linestyle='-',color='m')
plt.title("Elbow Method for Optional K")
plt.xlabel("Number of Clusters")
plt.ylabel("WCSS(Inertia)")
plt.show()
```

C:\Users\DELL\Downloads\priya python\lib\site-packages\sklearn\cluster_kmeans.py:1036: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
warnings.warn(



```
In [25]: kmeans=KMeans(n_clusters=3,random_state=42)
kmeans_labels=kmeans.fit_predict(x_scaled)

dbSCAN=DBSCAN(eps=0.5,min_samples=5)
dbSCAN_labels=dbSCAN.fit_predict(x_scaled)

kmeans_silhouette=silhouette_score(x_scaled,kmeans_labels)
dbSCAN_silhouette=silhouette_score(x_scaled[dbSCAN_labels != -1],dbSCAN_labels[dbSCAN_labels != -1])

print(f"silhouette_score for kmeans:{kmeans_silhouette:.3f}")
print(f"silhouette_score for dbSCAN:{dbSCAN_silhouette:.3f}")

silhouette_score for kmeans:0.459
silhouette_score for dbSCAN:0.653
```


```
In [55]: from sklearn.decomposition import PCA
pca=PCA(n_components=2)
x_reduced=pca.fit_transform(x_scaled)
fig,axs=plt.subplots(1,2,figsize=(10,4))

sns.scatterplot(x=x_reduced[:,0],y=x_reduced[:,1],hue=kmeans_labels,palette="dark:salmon_r",ax=axs[0])
axs[0].set_title("K-Means Clustering")
```

```

sns.scatterplot(x=x_reduced[:,0],y=x_reduced[:,1],hue=dbscan_labels,palette="coolwarm",ax=axs[1])
axs[1].set_title("DBSCAN Clustering")
plt.show()

```



```

In [56]: if "species" in df.columns:
    ground_truth=df["species"]
    print(f"Ground Truth Comparison:\n{pd.crosstab(ground_truth,kmeans_labels)}")
    print(f"\nGround Truth Comparison:\n{pd.crosstab(ground_truth[dbscan_labels!=-1],dbscan_labels)}")

Ground Truth Comparison:
K-Means Labels vs Species:
   col_0      0     1     2
species
0          0  50   0
1         11   0  39
2         36   0  14

Ground Truth Comparison:
DBSCAN Labels vs Species:
   col_0      0     1
species
0         44   0
1         0  39
2         0  32

```

In []:

In []: .

Program 12

```
In [ ]: !pip install mlxtend
```

```
In [13]: import pandas as pd
from mlxtend.frequent_patterns import apriori,association_rules
from mlxtend.preprocessing import TransactionEncoder

grocery_data=pd.read_csv(r"D:\PRIYA\SEM 05\dataset\groceries - groceries.csv")
transactions=grocery_data.iloc[:,1:].values.tolist()
transactions=[[item for item in transaction if isinstance(item,str)]
              for transaction in transactions]

te=TransactionEncoder()
te_ary=te.fit(transactions).transform(transactions)
one_hot_data=pd.DataFrame(te_ary,columns=te.columns_)

min_support=0.05
frequent_itemsets=apriori(one_hot_data,min_support=min_support,use_colnames=True)

min_confidence=0.1
rules=association_rules(frequent_itemsets,metric="confidence",min_threshold=min_confidence,num_itemsets=frequent_itemsets)

print("frequent itemsets:")
print(frequent_itemsets)
print("\n association rules:")
print(rules)

frequent itemsets:
 support           itemsets
0    0.052466          (beef)
1    0.080529      (bottled beer)
2    0.110524   (bottled water)
3    0.064870     (brown bread)
4    0.055414        (butter)
5    0.077682     (canned beer)
6    0.082766     (citrus fruit)
7    0.058058       (coffee)
8    0.053279       (curd)
9    0.063447   (domestic eggs)
10   0.058973      (frankfurter)
11   0.072293 (fruit/vegetable juice)
12   0.058566      (margarine)
13   0.052364      (napkins)
14   0.079817      (newspapers)
15   0.193493   (other vegetables)
16   0.088968       (pastry)
17   0.075648      (pip fruit)
18   0.057651       (pork)
19   0.183935      (rolls/buns)
20   0.108998   (root vegetables)
21   0.093950       (sausage)
22   0.098526      (shopping bags)
23   0.174377       (soda)
24   0.104931   (tropical fruit)
25   0.071683   (whipped/sour cream)
26   0.255516      (whole milk)
27   0.139502       (yogurt)
28   0.074835 (whole milk, other vegetables)
29   0.056634   (whole milk, rolls/buns)
30   0.056024      (whole milk, yogurt)

association rules:
 antecedents           consequents antecedent support \
0   (whole milk)   (other vegetables)  0.255516
1   (other vegetables)   (whole milk)  0.193493
2   (whole milk)   (rolls/buns)  0.255516
3   (rolls/buns)   (whole milk)  0.183935
4   (whole milk)       (yogurt)  0.255516
5       (yogurt)   (whole milk)  0.139502

 consequent support      support confidence      lift representativity \
0       0.193493  0.074835  0.292877  1.513634      1.0
1       0.255516  0.074835  0.386758  1.513634      1.0
2       0.183935  0.056634  0.221647  1.205032      1.0
3       0.255516  0.056634  0.307905  1.205032      1.0
4       0.139502  0.056024  0.219260  1.571735      1.0
5       0.255516  0.056024  0.401603  1.571735      1.0

 leverage conviction  zhangs_metric jaccard certainty kulczynski
0  0.025394    1.140548     0.455803  0.200000  0.123228  0.339817
1  0.025394    1.214013     0.420750  0.200000  0.176286  0.339817
2  0.009636    1.048452     0.228543  0.147942  0.046213  0.264776
3  0.009636    1.075696     0.208496  0.147942  0.070369  0.264776
4  0.020379    1.102157     0.488608  0.165267  0.092688  0.310432
5  0.020379    1.244132     0.422732  0.165267  0.196226  0.310432
```